

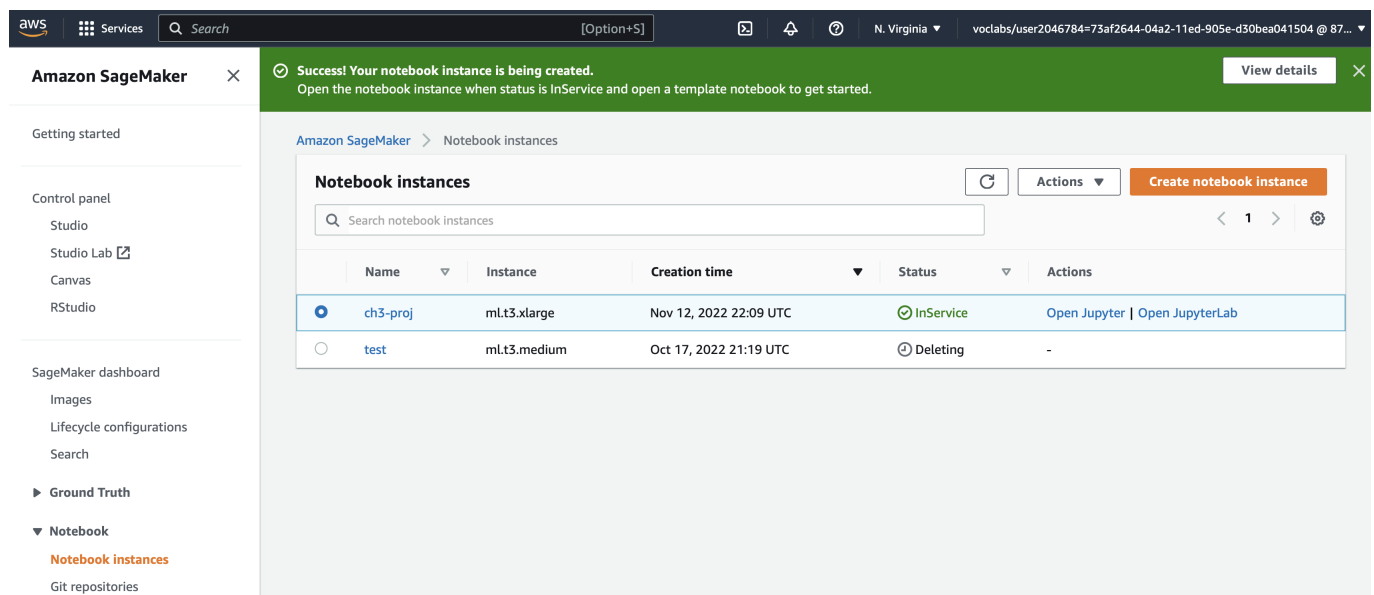
udacity-aws-machine-learning-engineer-nanodegree-ch3-operationalizing-an-aws-ml-project

udacity-aws-machine-learning-engineer-nanodegree-ch3-operationalizing-an-aws-ml-project

Initial setup, training and deployment

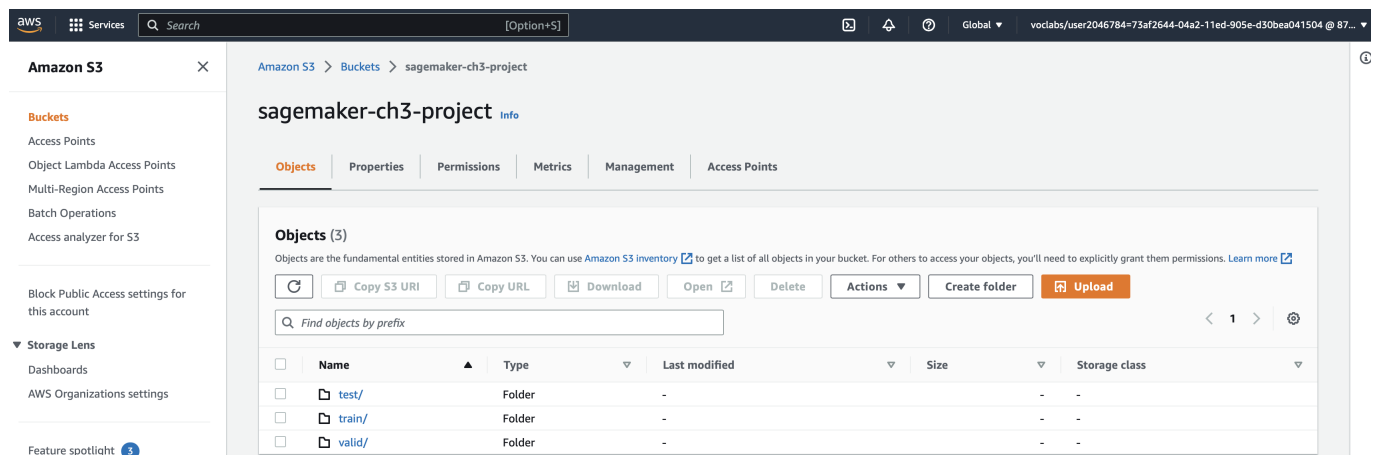
Initial Setup and Justification

I chose an ml.t3.xlarge, because I encountered a memory insufficiency error using ml.t3.medium (to save costs). Per docs <https://docs.aws.amazon.com/sagemaker/latest/dg/howitworks-create-ws.html>, if more memory is needed, recommend large or xlarge.



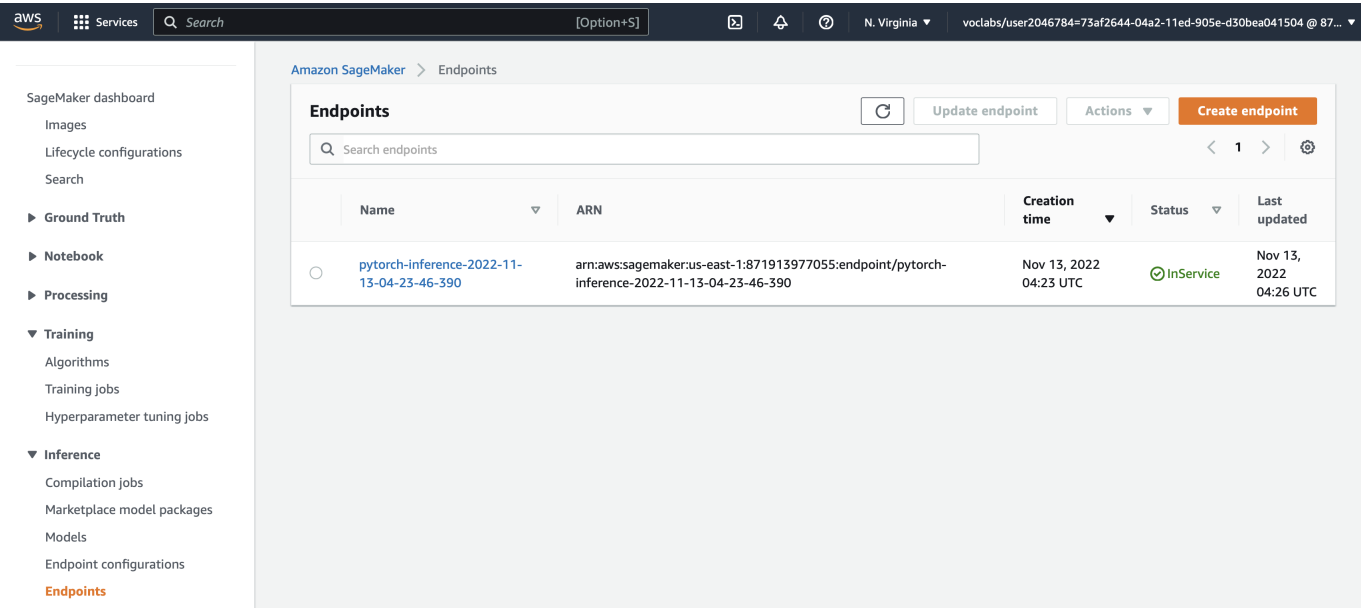
After you set up an S3 bucket, take a screenshot showing that you've set up an S3 bucket. Include this screenshot in your final submission.

Screenshot of S3 bucket



Training and Deployment

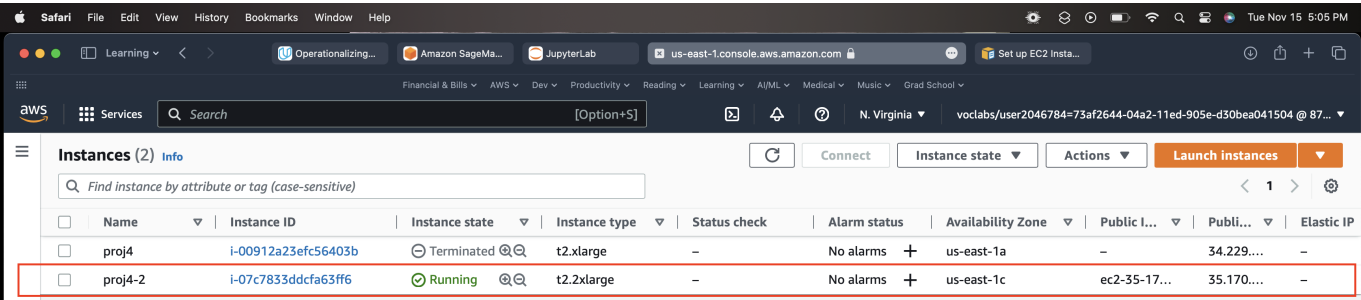
Screenshot of deployed endpoint in Inference > Endpoints



EC2 Training

Justification for EC2 Instance

For EC2 Training I decided to choose something relatively low cost, but with more memory than the base tier - t2.xlarge (4 vCPUs and 16 Gib memory) This is to strike a balance between cost and performance.



Saved EC2 Model

```

● ● ● Downloads — ec2-user@ip-172-31-4-113:~/TrainedModels — ssh -i proj4-2-key.pem ec2-user@ec2-34-200-214-195.compute-1.amazonaws.com — 98x30
(pytorch) [ec2-user@ip-172-31-4-113 ~]$ cd dogImages
(pytorch) [ec2-user@ip-172-31-4-113 dogImages]$ ls
test train valid
(pytorch) [ec2-user@ip-172-31-4-113 dogImages]$ cd ..
(pytorch) [ec2-user@ip-172-31-4-113 ~]$ python solution.py
/opt/conda/envs/pytorch/lib/python3.9/site-packages/torchvision/models/_utils.py:208: UserWarning:
The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use
'weights' instead.
  warnings.warn(
/opt/conda/envs/pytorch/lib/python3.9/site-packages/torchvision/models/_utils.py:223: UserWarning:
Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be r
emoved in the future. The current behavior is equivalent to passing `weights=ResNet50_Weights.IMAG
ENET1K_V1`. You can also use `weights=ResNet50_Weights.DEFAULT` to get the most up-to-date weights
.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /home/ec2-user/.cache/
torch/hub/checkpoints/resnet50-0676ba61.pth
100%|██████████████████████████████████████████████████████████████████████████████| 97.8M/97.8M [00:00<00:00, 114MB/s]
Starting Model Training
saved
(pytorch) [ec2-user@ip-172-31-4-113 ~]$ ls
BUILD_FROM_SOURCE_PACKAGES_LICENCES  LINUX_PACKAGES_LIST          THIRD_PARTY_SOURCE_CODE_URLS
dogImages                             nvidia-acknowledgements     TrainedModels
dogImages.zip                         PYTHON_PACKAGES_LICENSES
LINUX_PACKAGES_LICENSES              solution.py
(pytorch) [ec2-user@ip-172-31-4-113 ~]$ cd TrainedModels
(pytorch) [ec2-user@ip-172-31-4-113 TrainedModels]$ ls
model.pth
(pytorch) [ec2-user@ip-172-31-4-113 TrainedModels]$ █

```

Differences between SageMaker code and EC2 training code

1. Reduced use of logger.

```
logger.setLevel(logging.DEBUG)
logger.addHandler(logging.StreamHandler(sys.stdout))``
```

- ## 2. No main function or arg parser

Lambda Function

I created a lambda function using the provided starter file. Within the function, I provided the endpoint name.

Test inference

This is my test event:

Here are the results of the invocation:

4 / 6

```
-0.2843036949634552, -1.3399943113327026, 0.19547127187252045,  
-0.9253167510032654, 0.07587957382202148, 0.4271191358566284,  
-0.9859912991523743, 0.67887943983078, -1.3299245834350586,  
-0.6822599172592163, 0.29655855894088745, -0.5139930248260498,  
1.0835564136505127, -0.7982521057128906, 0.8554990291595459,  
0.362626850605011, 1.2404662370681763, -0.07426321506500244,  
0.10146191716194153, -0.68869549036026, -0.48106521368026733,  
-0.8763693571090698, -0.571552574634552, -0.6792696714401245,  
-0.5206713080406189, -1.9426512718200684, -2.419771671295166,  
-0.3339347839355469, -0.24408461153507233, -0.08373317122459412,  
0.8919371962547302, -2.856651782989502, -0.7000316977500916,  
0.15362446010112762, -0.7970098853111267, -0.35622841119766235,  
-1.9542192220687866, -2.7021138668060303, 0.01772058755159378,  
-1.431732416152954, -0.9516850113868713, -1.1090036630630493,  
-1.1546827554702759, -0.2241167426109314, 1.5300416946411133,  
-1.715854525566101, -0.03903423622250557, -1.1843053102493286,  
-1.9356435537338257, -0.12285503000020981, -1.4931055307388306,  
-0.15628209710121155, -0.536418616771698, -1.1524658203125,  
0.8359684944152832, -1.9632141590118408, -0.04063112288713455,  
0.36780649423599243, -1.129619836807251, -1.367497205734253,  
-0.0011828616261482239, -0.6425749063491821, -0.7356793880462646,  
1.0322339534759521, -0.8544481992721558, -1.4675880670547485,  
-1.728093147277832, -1.2840062379837036, -0.4094080626964569,  
-0.38233983516693115, -1.1601954698562622, 0.1524370163679123,  
-1.3115702867507935, -1.653199315071106, -2.022322654724121,  
-1.609832525253296, 0.21691259741783142, -2.072298526763916,  
-1.5164365768432617, -1.1745740175247192, -0.8381267189979553,  
-0.9472677707672119, -0.36459943652153015, -0.6767148971557617,  
-0.06865213811397552, 0.7558403015136719, -1.1166062355041504,  
-0.859592080116272, -1.0786547660827637, -0.19210000336170197,  
-1.1853835582733154, 0.21012145280838013, -1.9923133850097656,  
-0.22051474452018738, -1.0778053998947144, -1.6710342168807983,  
-1.1564878225326538, -1.1558247804641724, -1.6072654724121094,  
-1.6592135429382324, -1.4082069396972656, -1.3968602418899536,  
-1.9419918060302734, -1.5573835372924805, -2.394296646118164,  
-0.9858585596084595, -1.873814344406128]]"  
}
```

lam role

The screenshot shows the AWS IAM console interface. The left sidebar contains navigation links for Identity and Access Management (IAM), Access management, Access reports, and Related consoles. The main content area displays the details for the role 'proj4-inf-role-81cwp4b9'. The 'Summary' tab is active, showing the role's creation date (November 16, 2022, 19:41 UTC-07:00), last activity (None), ARN, and maximum session duration (1 hour). Below the summary, there are tabs for Permissions, Trust relationships, Tags, Access Advisor, and Revoke sessions. The 'Permissions' tab is selected, showing a list of permissions policies attached to the role. The list includes 'AWSLambdaBasicExecutionRole-fe8daa81-1bdd-4877-bc64-f95191fe...' (Customer managed) and 'AmazonSageMakerFullAccess' (AWS managed). The 'Permissions boundary' section indicates that no boundary is currently set.

Potential security vulnerabilities

I think my workspace is secure, but in general note that: roles that have "FullAccess" policies attached may be too permissive and may lead to problems. Roles that are old or inactive may lead to vulnerabilities because they may belong to people who are no longer working on the project and who may not be careful about ensuring the project's success.

Concurrency

I set up provisioned concurrency of 2, so that my lambda can handle more requests concurrently.

Autoscaling

I configured my endpoint to have autoscaling. I chose to have a scale in and scale out cool down period of 30 seconds so that the endpoint will have less delay before spinning up additional instances, as well as less delay before shutting down any excess instance. I chose a value of 10 for `SageMakerVariantInvocationsPerInstance`, which is my chosen level of how many invocations my endpoint instance should be able to handle before I want to spin up an additional instance and as a target to decide when to shut down excess instances.